

template-

November 17, 2023

1 Your Title Here

Name(s): (your name(s) here)

Website Link: https://lr580.github.io/power_outages_stats/

1.1 Code

```
[ ]: import pandas as pd
import numpy as np
import os
import plotly.express as px
pd.options.plotting.backend = 'plotly'
```

1.1.1 Cleaning and EDA

Data Cleaning

Read the Data Since the data is not large(as stated above, only over a thousand rows), we opened it beforehand via Microsoft Excel and found that the columns names are located at 6th rows, while the data starts at 8th row. And it's noticed that the first column is empty, which means that the real column starts at the second column. Therefore, we need to open the dataset starting from the 8th row and the 2nd column, while reading the 6th column as column names.

To realize this, we first use `header=5` to skip the first 5 rows and uses the 6th row as the names for columns. Then, we drop the first row(actually the 7th row in the raw file) and the first column.

```
[ ]: # skip first 5 rows
data = pd.read_excel("outage.xlsx", header=5)
# drop the 7th row
data = data.drop(0)
# drop the 1st column
data = data.drop(data.columns[0], axis=1)
# show the data
print(data)
```

| | OBS | YEAR | MONTH | U.S._STATE | POSTAL.CODE | NERC.REGION | \ |
|---|-----|--------|-------|------------|-------------|-------------|---|
| 1 | 1.0 | 2011.0 | 7.0 | Minnesota | MN | MRO | |
| 2 | 2.0 | 2014.0 | 5.0 | Minnesota | MN | MRO | |
| 3 | 3.0 | 2010.0 | 10.0 | Minnesota | MN | MRO | |

| | | | | | | |
|------|--------|--------|------|--------------|-----|------|
| 4 | 4.0 | 2012.0 | 6.0 | Minnesota | MN | MRO |
| 5 | 5.0 | 2015.0 | 7.0 | Minnesota | MN | MRO |
| ... | ... | ... | ... | ... | ... | ... |
| 1530 | 1530.0 | 2011.0 | 12.0 | North Dakota | ND | MRO |
| 1531 | 1531.0 | 2006.0 | NaN | North Dakota | ND | MRO |
| 1532 | 1532.0 | 2009.0 | 8.0 | South Dakota | SD | RFC |
| 1533 | 1533.0 | 2009.0 | 8.0 | South Dakota | SD | MRO |
| 1534 | 1534.0 | 2000.0 | NaN | Alaska | AK | ASCC |

| | CLIMATE.REGION | ANOMALY.LEVEL | CLIMATE.CATEGORY | OUTAGE.START.DATE | \ |
|------|--------------------|---------------|------------------|---------------------|-----|
| 1 | East North Central | -0.3 | normal | 2011-07-01 00:00:00 | |
| 2 | East North Central | -0.1 | normal | 2014-05-11 00:00:00 | |
| 3 | East North Central | -1.5 | cold | 2010-10-26 00:00:00 | |
| 4 | East North Central | -0.1 | normal | 2012-06-19 00:00:00 | |
| 5 | East North Central | 1.2 | warm | 2015-07-18 00:00:00 | |
| ... | ... | ... | ... | ... | ... |
| 1530 | West North Central | -0.9 | cold | 2011-12-06 00:00:00 | |
| 1531 | West North Central | NaN | NaN | NaN | |
| 1532 | West North Central | 0.5 | warm | 2009-08-29 00:00:00 | |
| 1533 | West North Central | 0.5 | warm | 2009-08-29 00:00:00 | |
| 1534 | NaN | NaN | NaN | NaN | |

| | POPCT_URBAN | POPCT_UC | POPDEN_URBAN | POPDEN_UC | POPDEN_RURAL | \ |
|------|-------------|----------|--------------|-----------|--------------|------|
| 1 | ... | 73.27 | 15.28 | 2279 | 1700.5 | 18.2 |
| 2 | ... | 73.27 | 15.28 | 2279 | 1700.5 | 18.2 |
| 3 | ... | 73.27 | 15.28 | 2279 | 1700.5 | 18.2 |
| 4 | ... | 73.27 | 15.28 | 2279 | 1700.5 | 18.2 |
| 5 | ... | 73.27 | 15.28 | 2279 | 1700.5 | 18.2 |
| ... | ... | ... | ... | ... | ... | ... |
| 1530 | ... | 59.9 | 19.9 | 2192.2 | 1868.2 | 3.9 |
| 1531 | ... | 59.9 | 19.9 | 2192.2 | 1868.2 | 3.9 |
| 1532 | ... | 56.65 | 26.73 | 2038.3 | 1905.4 | 4.7 |
| 1533 | ... | 56.65 | 26.73 | 2038.3 | 1905.4 | 4.7 |
| 1534 | ... | 66.02 | 21.56 | 1802.6 | 1276 | 0.4 |

| | AREAPCT_URBAN | AREAPCT_UC | PCT_LAND | PCT_WATER_TOT | PCT_WATER_INLAND |
|------|---------------|------------|-----------|---------------|------------------|
| 1 | 2.14 | 0.6 | 91.592666 | 8.407334 | 5.478743 |
| 2 | 2.14 | 0.6 | 91.592666 | 8.407334 | 5.478743 |
| 3 | 2.14 | 0.6 | 91.592666 | 8.407334 | 5.478743 |
| 4 | 2.14 | 0.6 | 91.592666 | 8.407334 | 5.478743 |
| 5 | 2.14 | 0.6 | 91.592666 | 8.407334 | 5.478743 |
| ... | ... | ... | ... | ... | ... |
| 1530 | 0.27 | 0.1 | 97.599649 | 2.401765 | 2.401765 |
| 1531 | 0.27 | 0.1 | 97.599649 | 2.401765 | 2.401765 |
| 1532 | 0.3 | 0.15 | 98.307744 | 1.692256 | 1.692256 |
| 1533 | 0.3 | 0.15 | 98.307744 | 1.692256 | 1.692256 |
| 1534 | 0.05 | 0.02 | 85.761154 | 14.238846 | 2.901182 |

[1534 rows x 56 columns]

Combine Date Column and Time Column It's not convenience if a specific datetime is described by two columns, so first we combine the `OUTAGE.START.DATE` column and the `OUTAGE.START.TIME` column into a single `OUTAGE.START` column with the `pd.Timestamp` type.

Before the conversion, we first check the two columns to see its types and raw data.

```
[ ]: print(data['OUTAGE.START.DATE'].dtype)
      print(data['OUTAGE.START.TIME'].dtype)
      print(data.loc[1, 'OUTAGE.START.DATE'])
      print(data.loc[1, 'OUTAGE.START.TIME'])
      print(data.loc[240, 'OUTAGE.START.DATE'])
      print(data.loc[240, 'OUTAGE.START.TIME'])
```

```
object
object
2011-07-01 00:00:00
17:00:00
nan
nan
```

It's noticed that some rows (like 240th row shown above) have NaN, so we should consider such special case when converting. The steps is shown below.

```
[ ]: # clone the raw data
      df = data.copy()

      # data type conversion
      df['OUTAGE.START.DATE'] = pd.to_datetime(df['OUTAGE.START.DATE'])
      df['OUTAGE.START.TIME'] = pd.to_timedelta(df['OUTAGE.START.TIME'].astype(str))

      # combine two columns into one column
      df['OUTAGE.START'] = df['OUTAGE.START.DATE'] + df['OUTAGE.START.TIME']
```

We check the result of the conversion.

```
[ ]: print(df['OUTAGE.START'].dtype)
      print(df.loc[1, 'OUTAGE.START'], df.loc[1, 'OUTAGE.START.TIME'], df.loc[1, 'OUTAGE.
      ↪ START.DATE'])
      print(df.loc[240, 'OUTAGE.START'], df.loc[240, 'OUTAGE.START.TIME'], df.
      ↪ loc[240, 'OUTAGE.START.DATE'])
```

```
datetime64[ns]
2011-07-01 17:00:00 0 days 17:00:00 2011-07-01 00:00:00
NaT NaT NaT
```

The cleaned dataframe's datatype for each useful column is listed below.

```
[ ]: selected_columns = ['OUTAGE.START', 'OUTAGE.DURATION', 'U.S._STATE', 'CLIMATE.
↳CATEGORY', 'OUTAGE.RESTORATION.DATE', 'CAUSE.CATEGORY.DETAIL']
print(df[selected_columns].dtypes)
df = df[selected_columns] # drop other useless columns
```

```
OUTAGE.START          datetime64[ns]
OUTAGE.DURATION        object
U.S._STATE            object
CLIMATE.CATEGORY       object
OUTAGE.RESTORATION.DATE object
CAUSE.CATEGORY.DETAIL  object
dtype: object
```

The cleaned dataframe is shown below(with only representative rows selected for display).

```
[ ]: selected_rows = df.loc[[5,4,3,54,833]]
print(selected_rows)
with open('table1.txt', 'w') as f:
    f.write(selected_rows.to_markdown(index=False))
```

| | OUTAGE.START | OUTAGE.DURATION | U.S._STATE | CLIMATE.CATEGORY | \ |
|-----|---------------------|-----------------|------------|------------------|---|
| 5 | 2015-07-18 02:00:00 | 1740 | Minnesota | warm | |
| 4 | 2012-06-19 04:30:00 | 2550 | Minnesota | normal | |
| 3 | 2010-10-26 20:00:00 | 3000 | Minnesota | cold | |
| 54 | 2014-01-24 00:00:00 | 108653 | Wisconsin | cold | |
| 833 | 2013-08-12 11:55:00 | 4 | Oregon | normal | |

| | OUTAGE.RESTORATION.DATE | CAUSE.CATEGORY.DETAIL | OUTAGE.RESTORATION.MISSING | \ |
|-----|-------------------------|-----------------------|----------------------------|---|
| 5 | 2015-07-19 00:00:00 | NaN | False | |
| 4 | 2012-06-20 00:00:00 | thunderstorm | False | |
| 3 | 2010-10-28 00:00:00 | heavy wind | False | |
| 54 | 2014-04-09 00:00:00 | Coal | False | |
| 833 | 2013-08-12 00:00:00 | suspicious activity | False | |

| | OUTAGE.DURATION.MISSING | OUTAGE.START.MISSING | \ |
|-----|-------------------------|----------------------|---|
| 5 | False | False | |
| 4 | False | False | |
| 3 | False | False | |
| 54 | False | False | |
| 833 | False | False | |

| | CAUSE.CATEGORY.DETAIL.MISSING | RANDOM | NO_RANDOM |
|-----|-------------------------------|--------|-----------|
| 5 | True | a | False |
| 4 | False | b | False |
| 3 | False | b | True |
| 54 | False | b | True |
| 833 | False | b | False |

Exploratory Data Analysis

Univariate Analysis In the univariate analysis, we would analyze the distribution of the outage duration and U.S. states.

Distribution of Outage Duration First, we write codes to analyse our first chosen column: outage duration.

```
[ ]: fig = px.histogram(df, x='OUTAGE.DURATION', title = 'Distribution of Outage_
    ↳Duration')
fig.update_layout(xaxis_title='Outage Duration (Minutes)')
fig.show()
```

This shows that exception from very few data, most of the duration is short. About half outages are less than 1,000 minutes, and most are less than 5,00 minutes, while only few outages lasting longer than that. Also, it means that the variance may be large since the minimal and maximal values differ a lot.

For convenience, we write a helper function to export the plotly figure into HTML file.

```
[ ]: def export_plotly_fig(fig, filename):
    fig.write_html(filename, include_plotlyjs='cdn')
```

We export the fig using the above function to local disk, which will be useful later in next part.

```
[ ]: export_plotly_fig(fig, 'univariate1.html')
```

Distribution of Outage U.S. States Then, we use the same way to analyze the distribution of U.S. states of the outages.

```
[ ]: # sorted by counts
count_df = df['U.S._STATE'].value_counts()

fig = px.bar(count_df, y='U.S._STATE', title = 'Distribution of Outage\'s U.S._
    ↳State')
fig.update_layout(xaxis_title='Outage U.S. State', yaxis_title='counts')
fig.show()
export_plotly_fig(fig, 'univariate2.html')
```

This shows that some states frequently occur power outages(more than 50 times in total), and most others are less frequent, even several states only have several power outages. It's seems that the distribution is similar to gaussian distribution.

Bivariate Analysis Then, we do bivariate analysis between the outage duration and the state where the outage occurs.

```
[ ]: fig = px.scatter(df, x = 'U.S._STATE', y = 'OUTAGE.DURATION', title =_
    ↳'Distribution of Outage Duration between Different States')
fig.update_layout(xaxis_title='U.S. State', yaxis_title='Outage Duration')
fig.show()
export_plotly_fig(fig, 'bivariate1.html')
```

The scatter plot shows that they may exist very strong correlation between the outage duration and the state. We could say that there may have positive relationship between the two variables.

To further explore it, we calculate the average outage duration minutes of different states and plot a line graph.

```
[ ]: # calculate the average duration for each states
avg_df = df.groupby('U.S._STATE')['OUTAGE.DURATION'].mean()
# sorted in ascending order
avg_df = avg_df.sort_values()
fig = px.line(avg_df, title = 'Average Outage Duration of Different States')
fig.update_layout(xaxis_title='U.S. State')
fig.show()
export_plotly_fig(fig, 'bivariate2.html')
```

It seems that the average outage duration is different between different states. To check it further, we calculate the median, and the quartile below.

```
[ ]: df2 = df.copy()
df2.dropna(subset=['OUTAGE.DURATION'])
df2['OUTAGE.DURATION'] = df2['OUTAGE.DURATION'].astype(float)
stat_df = df2.groupby('U.S._STATE')['OUTAGE.DURATION'].agg(['mean', 'median'])
stat_df = stat_df.sort_values(by='mean')
q1 = df2.groupby('U.S._STATE')['OUTAGE.DURATION'].quantile(0.25)
q3 = df2.groupby('U.S._STATE')['OUTAGE.DURATION'].quantile(0.75)
stat_df = pd.merge(stat_df, q1, on='U.S._STATE')
stat_df = pd.merge(stat_df, q3, on='U.S._STATE')
stat_df.columns = ['mean', 'median', 'Q1', 'Q3']
fig = px.line(stat_df, title = 'Statistics Outage Duration of Different States')
fig.update_layout(xaxis_title='U.S. State')
fig.show()
export_plotly_fig(fig, 'bivariate3.html')
```

It seems that the duration time differs a lot between different states. So we can roughly conclude that it may have strong correlation.

Interesting Aggregates In the above mentioned bivariate analysis, we get a grouped table `stat_df`. Here, we show it again in table form. Below shows some rows of a grouped table aggregated by states and calculated the mean, median, and quartile of the outage duration.

```
[ ]: print(stat_df.loc[['Mississippi', 'Hawaii', 'Massachusetts', 'Tennessee', 'New_York']])
with open('table2.txt', 'w') as f:
    f.write(stat_df.loc[['Mississippi', 'Hawaii', 'Massachusetts', 'Tennessee', 'New_York']].to_markdown())
```

| | mean | median | Q1 | Q3 |
|-------------|------------|--------|-------|---------|
| U.S._STATE | | | | |
| Mississippi | 84.000000 | 17.5 | 4.0 | 97.50 |
| Hawaii | 845.400000 | 543.0 | 237.0 | 1367.00 |

| | | | | |
|---------------|-------------|--------|-------|---------|
| Massachusetts | 944.166667 | 211.0 | 19.0 | 1443.75 |
| Tennessee | 1041.967742 | 310.0 | 39.0 | 1230.00 |
| New York | 6034.957143 | 2880.0 | 268.5 | 8156.25 |

Furthermore, we want to explore the relationship between different climate category (i.e. CLIMATE.CATEGORY column) and the outage duration of different states. So we plot an pivot table below (Also, only show some rows).

```
[ ]: pivot_table = df2.pivot_table(index='U.S._STATE', columns='CLIMATE.CATEGORY',
    ↪values='OUTAGE.DURATION', aggfunc='mean')
rows = pivot_table.loc[['Wisconsin', 'Hawaii', 'Massachusetts', 'Tennessee',
    ↪'New York']]
print(rows)
with open('table3.txt', 'w') as f:
    f.write(rows.to_markdown())
```

| CLIMATE.CATEGORY | cold | normal | warm |
|------------------|--------------|-------------|-------------|
| U.S._STATE | | | |
| Wisconsin | 12545.555556 | 3962.555556 | 1605.000000 |
| Hawaii | 1367.000000 | 205.500000 | 1224.500000 |
| Massachusetts | 160.333333 | 1159.384615 | 721.000000 |
| Tennessee | 1476.583333 | 1015.750000 | 341.857143 |
| New York | 8914.576923 | 3673.562500 | 6092.833333 |

To observe it more intuitively, we plot the table below.

```
[ ]: fig = px.bar(pivot_table, title = 'Average Outage Duration of Different Climate_
    ↪Category of Different States')
fig.update_layout(xaxis_title='U.S. State', yaxis_title='Average Outage_
    ↪Duration')
fig.show()
export_plotly_fig(fig, filename='bar_climate_category.html')
```

It seems that the outage duration differ a lot in different climate. The warmer, the longer it lasts. But it's not necessary that it's true, so we'd check that in the hypothesis test below.

1.1.2 Assessment of Missingness

NMAR Analysis One of the column in our dataset with missing values that is possibly NMAR is the OUTAGE.DURATION column. The reason for its missingness is most likely due to some reasons that why the OUTAGE.START.DATE or more frequently, the OUTAGE.RESTORATION.DATE is missing. Any of these columns miss will lead to the OUTAGE.DURATION miss. However, there's no clear reasoning why they miss recorded in the dataset. If we want additional data to record the reason why the start and restoration datetime of a outage is missing, then we could make it MAR.

Missingness Dependency We constructed permutation tests to determine the relationship. We've decided to test the dependency between the missingness of OUTAGE.DURATION with two columns: OUTAGE.RESTORATION.DATE and CAUSE.CATEGORY.DETAIL.

OUTAGE.DURATION and OUTAGE.RESTORATION.DATE (MAR) To make our code suitable for more different columns, we first build a universal function to check whether the missingness of one column depends on another column.

It's no doubt that the missingness of the `OUTAGE.DURATION` depend on the missingness `OUTAGE.RESTORATION.DATE`. We can easily figure out this from both observing the data and reasoning in the reality. And now, we'd like check it by performing permutation tests.

More specifically, it means that the missingness of `OUTAGE.RESTORATION.DATE` affects the missingness of `OUTAGE.DURATION`. So, we first create an additional column `OUTAGE.RESTORATION.MISSING` to record the whether the `OUTAGE.RESTORATION.DATE` is missing.

```
[ ]: df['OUTAGE.RESTORATION.MISSING'] = df['OUTAGE.RESTORATION.DATE'].isna()
      print(df['OUTAGE.RESTORATION.MISSING'].tail())
```

```
1530    False
1531     True
1532    False
1533    False
1534     True
Name: OUTAGE.RESTORATION.MISSING, dtype: bool
```

Since it's a category column, we use TVD(Total Variation Distance) to perform permutation test. Recall that:

$$TVD(X, Y) = \frac{1}{2} \sum_{i=1}^n |X_i - Y_i|$$

The null hypothesis for the permutation test is that the specific column does not depend on another column. So if p-value is less than the significance level, we reject the null hypothesis, which means that we may think the column depend on another. Otherwise, we fail to reject the null hypothesis, which means that it's more possible to believe the column does not depend on another.

Below, we make codes to compute the TVD, run the permutation tests of 500 rounds, set the significance level to 0.05, calculate the p-value, draw the plot and get the conclusion.

```
[ ]: def TVD(df, col1, col2):
      pivoted = (
          df
          .pivot_table(index=col1, columns=col2, aggfunc='size')
          .apply(lambda x: x / x.sum() ).fillna(0)
      )
      return pivoted.diff(axis=1).iloc[:, -1].abs().sum() / 2

def report_perm(original_tvd, p_value, tvds, col1, col2):
    # plot
    fig = px.histogram(pd.DataFrame(tvds), x=0, nbins=20,
    ↪ histnorm='probability')
    fig.add_vline(x=original_tvd, line_color='red')
```



```

fig.add_annotation(text=f'<span style="color:red">Observed TVD =␣
↳{round(original_tvd, 2)}, p_value = {round(p_value ,2)}</span>',
                  x= 0.4, showarrow=False, y=0.1)
fig.update_layout(title = f"Permutation Test whether <br> {col1} Depends on␣
↳<br> {col2}", xaxis_title="TVD")
fig.show()
export_plotly_fig(fig, f'permutation_test_{col1}_{col2}.html')

# make conclusion
if p_value < 0.05:
    print(f'Reject the null hypothesis, \nit\'s more likely that {col1}␣
↳depends on {col2}.')
else:
    print(f'Fail to reject the null hypothesis, \nit\'s more likely that␣
↳{col1} does not depend on {col2}.')

def permutation_test(df, col1, col2, n_permutations=500):
    original_tvd = TVD(df, col1, col2)

    permuted_tvds = []
    shuffled = df.copy()
    for _ in range(n_permutations):
        shuffled[col1] = np.random.permutation(df[col1])
        permuted_tvds.append(TVD(shuffled, col1, col2))

    p_value = np.mean([tvd >= original_tvd for tvd in permuted_tvds])
    report_perm(original_tvd, p_value, permuted_tvds, col1, col2)
    return original_tvd, p_value, permuted_tvds

```

We use the codes to check whether the OUTAGE.DURATION missingness depends on the OUTAGE.RESTORATION missingness.

```

[ ]: df['OUTAGE.DURATION.MISSING'] = df['OUTAGE.DURATION'].isna()
tvd, p, tvds = permutation_test(df, 'OUTAGE.DURATION.MISSING', 'OUTAGE.
↳RESTORATION.MISSING')
print(tvd, p)

```

```

Reject the null hypothesis,
it's more likely that OUTAGE.DURATION.MISSING depends on
OUTAGE.RESTORATION.MISSING.
1.0 0.0

```

The result shows that p-value is 0.00, since the significance level is 5%, we reject the null hypothesis. Therefore, we conclude that it is highly possible that the missingness of the OUTAGE.DURATION depends on the missingness of the OUTAGE.RESTORATION.

We may get the same conclusion that the missingness of the OUTAGE.DURATION is also depends on the missingness of the OUTAGE.START, but the TVDs of the permutation tests should be larger than the previous ones, since the OUTAGE.START data miss less, which can be easily seen from the

dataset. The result below shows that it's rational.

```
[ ]: df['OUTAGE.START.MISSING'] = df['OUTAGE.START'].isna()
      tvd, p, tvds = permutation_test(df, 'OUTAGE.DURATION.MISSING', 'OUTAGE.START.
      ↳MISSING')
      print(tvd, p)
```

Reject the null hypothesis,
it's more likely that OUTAGE.DURATION.MISSING depends on OUTAGE.START.MISSING.
0.9678688524590164 0.0

We can easily find that, on the other hand, the missingness of the CAUSE.CATEGORY.DETAIL, have nothing to do with the missingness of OUTAGE.DURATION. Since it only depicts some additional information of the CAUSE.CATEGORY. So we should find that the missingness of the CAUSE.CATEGORY.DETAIL is independent of the missingness of the CAUSE.CATEGORY.DETAIL. Below, we will check our conjecture.

```
[ ]: df['CAUSE.CATEGORY.DETAIL.MISSING'] = df['CAUSE.CATEGORY.DETAIL'].isna()
      tvd, p, tvds = permutation_test(df, 'OUTAGE.DURATION.MISSING', 'CAUSE.CATEGORY.
      ↳DETAIL')
      print(tvd, p)
```

Fail to reject the null hypothesis,
it's more likely that OUTAGE.DURATION.MISSING does not depend on
CAUSE.CATEGORY.DETAIL.
0.009900990099009906 0.89

The plot result shows that p-value is 0.89, which is largely greater than the significance level of 5%, so we fail to reject the null hypothesis. Hence, we conclude that the missingness of OUTAGE.DURATION does not depend on the missingness of CAUSE.CATEGORY.DETAIL.

1.1.3 Hypothesis Testing

1.1.4 Permutation Test

Going back to our investigation topic, we are investigating does climate category affect the duration of a power outage. Recalling the plot in the interesting aggregates part, we intuitively propose the idea that there are some relationships between the climate and the duration of power outages. But observation alone cannot be a good indicator. So we determined a good way to test is making a hypothesis test.

We think that permutation test on the distribution of outage durations in the warm climate areas and the distribution in the cold climate areas to see whether there is an actual increase in duration time to warmer areas or not.

Null hypothesis: the outage duration time in the warm climate and the cold climate comes from the same distribution.

Alternative hypothesis: the outage duration time in the warm climate and the cold climate comes from different distributions.

Significance level: 5%.

Number of permutation test rounds: 500.

Test statistics: the outage duration is numeric data rather than categorical. And recall that numeric data uses diff of means, which the categorical data uses TVD. Thereby, we use the diff of means as our test statistics.

The value is numeric rather than category, so we use the absolute difference between the mean of the two groups to check it. Recall the formula:

$$\text{diff_of_means}(A, B) = \left| \frac{1}{n} \sum_{i=1}^n A_i - \frac{1}{m} \sum_{i=1}^m B_i \right|$$

Using the similar way, we write codes of the permutation test below.

```
[ ]: def diff_of_means(df, col1, col2, v1, v2):
    mean1 = df[df[col2] == v1][col1].mean()
    mean2 = df[df[col2] == v2][col1].mean()
    return abs(mean1 - mean2)
def report_hyp(observed, p_value, simulated, col1, col2, v1, v2):
    # plot
    fig = px.histogram(pd.DataFrame(simulated), x=0, nbins=20,
    histnorm='probability')
    fig.add_vline(x=observed, line_color='red')
    fig.add_annotation(text=f'<span style="color:red">Observed =
    round(observed, 2)}, p_value = {round(p_value, 2)}</span>',
    x= 0.4, showarrow=False, y=0.1)
    fig.update_layout(title = f"Empirical Distribution to check whether <br>
    {col1} Have Different Distribution when <br> {col2} is {v1} and {v2}",
    axis_title="Diff of Means")
    fig.show()
    export_plotly_fig(fig, filename=f'hyp_{col1}_{col2}_{v1}_{v2}.html')

    # make conclusion
    if p_value < 0.05:
        print(f'Reject the null hypothesis, \nit\'s more likely that {col1} is
    different when {col2} is {v1} and {v2}.')
    else:
        print(f'Fail to reject the null hypothesis, \nit\'s more likely that
    {col1} is same when {col2} is {v1} and {v2}.')
def hypothesis_test(df, col1, col2, v1, v2, rounds=500):
    observed = diff_of_means(df, col1, col2, v1, v2)
    df2 = df.copy()
    simulated = np.zeros(rounds)
    for _ in range(rounds):
        df2[col2] = df[col2].sample(frac=1).reset_index(drop=True)
        simulated[_] = diff_of_means(df2, col1, col2, v1, v2)
    p_value = np.mean(simulated >= observed)
    report_hyp(observed, p_value, simulated, col1, col2, v1, v2)
```

```
return observed, p_value
```

```
[ ]: hypothesis_test(df, 'OUTAGE.DURATION', 'CLIMATE.CATEGORY', 'warm', 'cold')
```

Fail to reject the null hypothesis,
it's more likely that OUTAGE.DURATION is same when CLIMATE.CATEGORY is warm and cold.

```
[ ]: (160.36121774568983, 0.712)
```

1.1.5 Conclusion

Out of our expectation, the result shows that p-value is 0.71, which is larger than 0.05, so we fail to reject the null hypothesis, which means that the outage duration time in the warm climate and the cold climate comes from the same distribution. That is, the outage duration have no relationship with the climate.

To check the correctness, we make a totally random column to check our correctness, it should be a same distribution below.

```
[ ]: df['RANDOM'] = np.random.choice(['a', 'b'], df.shape[0], p=[0.4, 0.6])  
hypothesis_test(df, 'OUTAGE.DURATION', 'RANDOM', 'a', 'b')
```

Fail to reject the null hypothesis,
it's more likely that OUTAGE.DURATION is same when RANDOM is a and b.

```
[ ]: (261.0528683297448, 0.448)
```

And the below should come from different distribution. We divide duration into two groups by whether it is greater than it's mean value. So obviously it's different distribution.

```
[ ]: mean = df['OUTAGE.DURATION'].mean()  
df['NO_RANDOM'] = df['OUTAGE.DURATION'] > mean  
hypothesis_test(df, 'OUTAGE.DURATION', 'NO_RANDOM', True, False)
```

Reject the null hypothesis,
it's more likely that OUTAGE.DURATION is different when NO_RANDOM is True and False.

```
[ ]: (7224.882703670063, 0.0)
```

Similarly, we can use the same function to test whether different states affect the distribution.

```
[ ]: hypothesis_test(df, 'OUTAGE.DURATION', 'U.S._STATE', 'New York', 'Tennessee')
```

Reject the null hypothesis,
it's more likely that OUTAGE.DURATION is different when U.S._STATE is New York and Tennessee.

```
[ ]: (4992.989400921659, 0.0)
```

It shows that it's highly possible that the outage duration differs a lot in the two states.